

IN-PERSON

SATURDAY | OCT 26 | 1400 AM - 1700 PM (ICT)

MongoDB Mini Meetup: MongoDB Data Modeling and Query Optimization



Piti Champeethong

Senior Consulting Engineer
MongoDB User Group Leader

Agenda

- Core Topologies
 - Replication
 - Sharding
- Core Data Modeling
 - Computed pattern
 - Inheritance pattern
 - Extended reference pattern
 - Schema versioning pattern
 - Subset pattern
 - Bucket pattern



Preparation

- <https://www.mongodb.com/docs/atlas/cli/current/install-atlas-cli/>
- `curl https://atlas-education.s3.amazonaws.com/sampledata.archive -o sampledata.archive`
- <https://www.mongodb.com/try/download/compass>




MongoDB – Document database

- Advantages
 - Full cloud-based developer data platform (Atlas)
 - Flexible document schemas
 - Widely supported and code-native data access
 - Change-friendly design
 - Powerful querying and analytics
 - Easy horizontal scale-out with sharding
 - Simple installation
- Disadvantages
 - Easy to get mistake by misunderstanding.



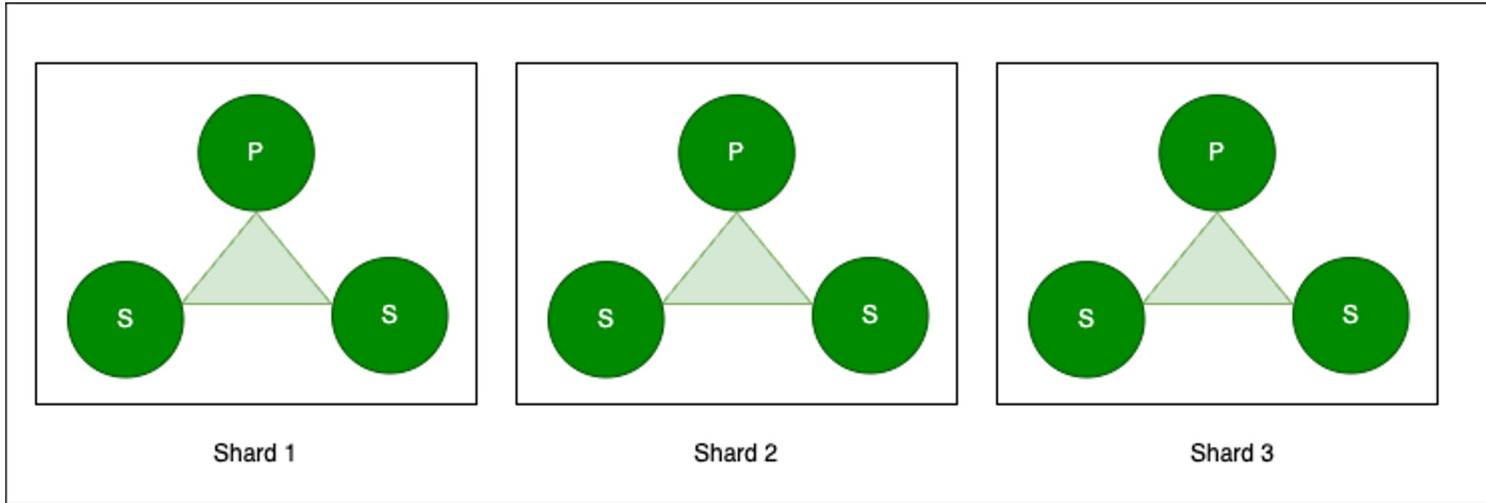
Replication



- Only primary node can accept write operations
- All nodes can accept read operations
- Maximum 7 nodes are vote member nodes.
- High Availability (HA)



Sharding



- Horizontal Scaling
- Collection sharding



Computed Pattern

```
{
  "_id": 1,
  "product_name": "Laptop",
  "reviews": [
    {
      "user": "Alice",
      "comment": "Great laptop",
      "stars": 5,
      "date": { "$date": "2021-01-01T12:30:00Z" }
    },
    {
      "user": "Bob",
      "comment": "Good laptop",
      "stars": 4,
      "date": { "$date": "2021-02-01T12:30:00Z" }
    }
  ],
  "computed_avg_rating": 4.5,
  "total_reviews": 2
}
```



Computed Pattern

Problem

- Documents are more similar than different
- Need to query the documents on their similitudes.

Use cases

- Internet of Things
- Event sourcing
- E-commerce



Inheritance Pattern

```
{
  "_id": 1,
  "type": "parent",
  "name": "John Doe",
  "age": 40
},
{
  "_id": 2,
  "type": "child",
  "name": "Jane Doe",
  "age": 10,
  "parent": 1
}
```

```
{
  "_id": 1,
  "patient_id": "P0001",
  "type": "surgical",
  "surgery_name": "Appendectomy",
  "surgeon": "Dr. Smith",
  "date": { "$date": "2024-01-01T10:00:00Z" }
},
{
  "_id": 2,
  "patient_id": "P0001",
  "type": "dental",
  "dentist": "Appendectomy",
  "procedure": "Tooth Extraction",
  "date": { "$date": "2024-02-01T10:00:00Z" }
}
```



Inheritance Pattern

Problem

- Documents are more similar than different
- Need to query the documents on their similitudes.

Use cases

- Single View
- Product Catalog
- Content Management
- Mobile Application



Extended Reference Pattern

```
{
  "_id": {
    "$oid": "66027d48d57fb8ef29b4bb33"
  },
  "product_name": "Laptop ABC",
  "price": 34000,
  "specs": {
    "processor": "Intel Core i3 7th Gen",
    "ram": "4 GB DDR4",
    "storage": "1 TB HDD",
    "os": "Windows 10 Home"
  }
}
```

products collection

```
{
  "_id": { "$oid": "66027daed57fb8ef29b4bb34" },
  "customer_id": "C7990",
  "order_date": { "$date": "2024-02-01T10:30:00.000Z" },
  "products": [
    {
      "product_id": {
        "$oid": "66027d48d57fb8ef29b4bb33"
      },
      "product_name": "Laptop ABC",
      "price": 34000,
      "quantity": 1
    }
  ]
}
```

orders collection



Extended Reference Pattern

Problem

- Too many joins in read operations
- Embedding leads to documents that are too big (16MB)

Use cases

- Catalog.
- Real-time analytics.
- Mobile Application.
- E-commerce



Schema Versioning Pattern

```
{
  "_id": {
    "$oid": "66027d48d57fb8ef29b4bb33"
  },
  "username": "Solo",
  "email": "solo@company.mail",
  "schema_version": 1
}
```

```
{
  "_id": {
    "$oid": "66027daed57fb8ef29b4bb34"
  },
  "username": "Sanji",
  "email": "sanji@company.mail",
  "token": "66027daed57fb8ef29b4bb34",
  "token_expired": {
    "$date": "2029-02-01T10:30:00.000Z"
  },
  "schema_version": 2
}
```



Schema Versioning Pattern

Problem

- Doing a schema migration without downtime.

Use cases

- Any application that can't sustain any downtimes



Subset Pattern

```
{
  "_id": {
    "$oid": "66027d48d57fb8ef29b4bb33"
  },
  "title": "MongoDB 101",
  "author": "jojo hakusho",
  "subset_reviews": [
    { "rev_id": 1, "user": "Bob", "rating": 10 },
    { "rev_id": 2, "user": "Alice", "rating": 9 },
    { "rev_id": 3, "user": "Mike", "rating": 8 },
    { "rev_id": 4, "user": "Koi", "rating": 8 },
    { "rev_id": 5, "user": "Henry", "rating": 8 }
  ]
}
```

books collection

```
[
  { "rev_id": 6, "user": "Mee", "rating": 10 },
  { "rev_id": 7, "user": "Maew", "rating": 9 },
  { "rev_id": 8, "user": "Kai", "rating": 8 },
  { "rev_id": 9, "user": "Pop", "rating": 8 },
  { "rev_id": 10, "user": "Pae", "rating": 8 },
  { "rev_id": 11, "user": "Bow", "rating": 10 },
  { "rev_id": 12, "user": "Row", "rating": 9 },
  { "rev_id": 13, "user": "Mai", "rating": 8 },
  { "rev_id": 14, "user": "Maa", "rating": 8 },
  { "rev_id": 15, "user": "Leo", "rating": 8 }
]
```

reviews collection



Subset Pattern

Problem

- Large documents are taking up a lot of space in memory.

Use cases

- List of reviews.
- List of comments.
- A long list of nearly anything kept in an array



Bucket Pattern

```
{
  "_id": {
    "$oid": "66027d48d57fb8ef29b4bb33"
  },
  "sensor_id": "S123",
  "date": 20240203,
  "readings": [
    { "value": 22, "ts": { "$date": "2024-02-03T08:00:01Z" } },
    { "value": 20, "ts": { "$date": "2024-02-03T08:05:01Z" } },
    { "value": 18, "ts": { "$date": "2024-02-03T08:10:00Z" } },
    { "value": 23, "ts": { "$date": "2024-02-03T08:15:01Z" } },
    { "value": 22, "ts": { "$date": "2024-02-03T08:20:01Z" } },
  ],
  "count": 5
}
```



Bucket Pattern

Problem

- Avoiding too many documents or documents too big.
- A one-to-many relationship that can't be embedded.

Use cases

- Internet of Things.
- Data Warehouse.
- One-to-many relationships with high cardinality. (A large number of different values)





(E)quality

(S)ort

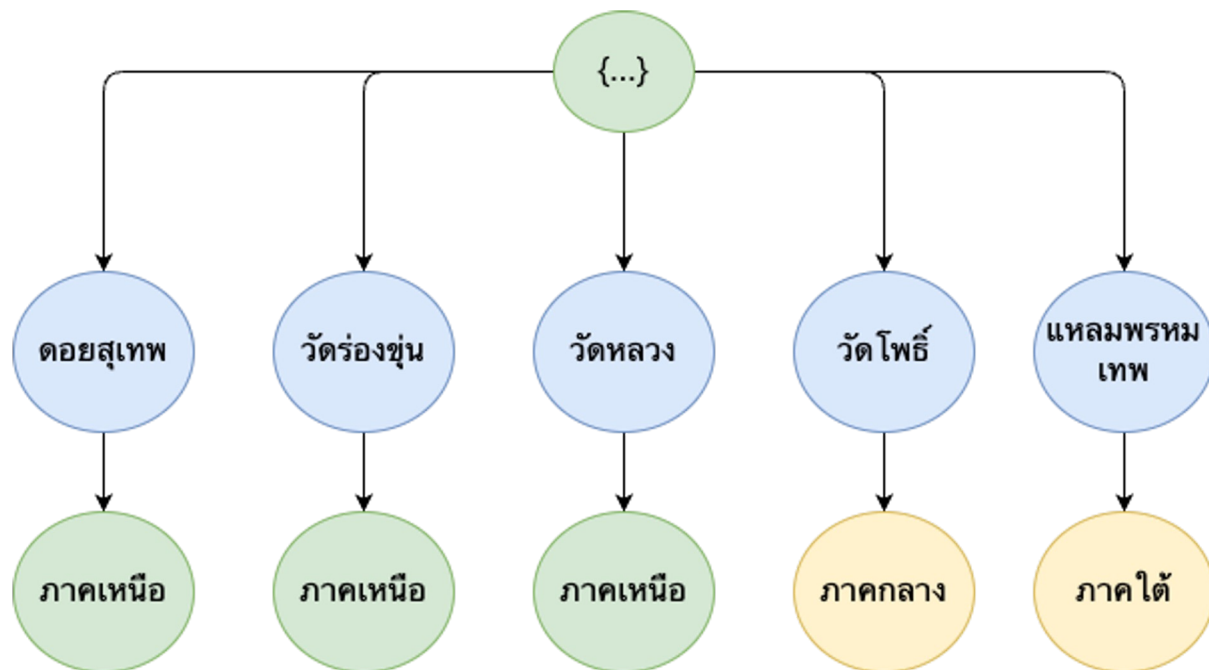
(R)ange

E-S-R

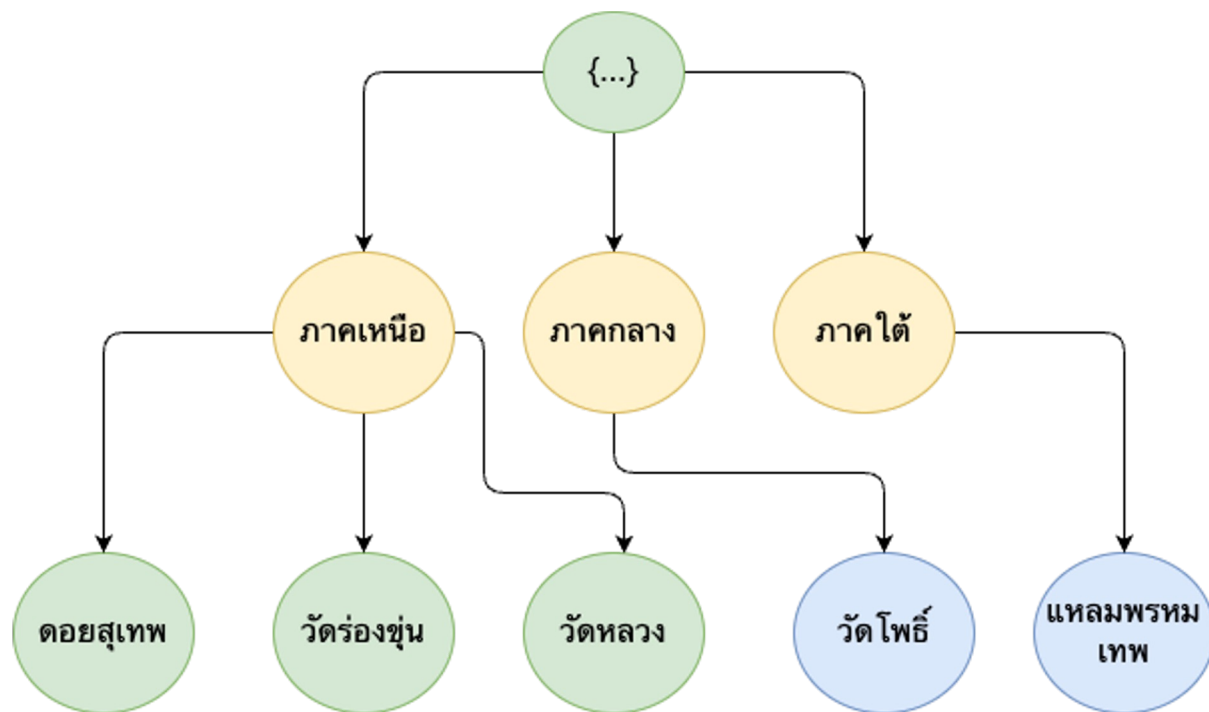
- Equality before Sort
- Equality before Range
- Sort before Range




E-S-R - Good




E-S-R - Better




E before R - Good




E before R - Better






S before R - Good



S before R - Better



Book recommendations



References

- <https://www.geopits.com/blog/mongodb-data-modeling-design-patterns.html>
- <https://www.mongodb.com/blog/post/new-data-modeling-learning-path-certification>
- <https://github.com/mongodbthailand/thmug-esr>
- <https://www.mongodb.com/developer/books>

